

作法

1. 從training data中把「同樣USER-ID」評書的紀錄收集起來
 - $U_i = \{r_1, r_2, \dots, r_j\}$
2. 計算各種統計量（平均、最大最小、偏態、峰態、10分位數）
3. 把 U_i 序列壓扁當成該User的特徵
4. 沒有資料的User特徵用中位數填入
 - 可能看過零本書、都是implicit rating、users.csv根本沒描述、...等
5. （忽略ISBN）
6. train lightGBM
 - 參數全部用預設
 - 微軟出的gradient boost decision tree套件，anaconda沒有要自己裝
7. 輸出track2預測、四捨五入到整數位當成track1預測

結論

1. 大部分使用者只評1本書（之後另寫一篇）
 - 沒算變異數是因為scipy呼叫了長度1的序列會掛掉XD
2. lightGBM train得比sklearn.ensemble.GradientBoostingClassifier快得多
 - 有很多錯誤函數可以選，之後再研究
3. lb score看來使用者評分的"scale"是有點重要（專門評低分的人vs都給8,9,10分的人）

可能的改進

1. 調lightGBM的參數
2. 對ISBN做簡單分析
3. 對user資料做一些視覺化（i.e. U_i ）

DataLoader / Imports

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import scipy
4 from learning2read.utils import DataLoader
5 def Data(key,**kwargs):
6     return DataLoader(r"/Users/qtwu/Downloads/data").load(key,**kwargs)
```

In [2]:

```
1 raw_user=Data("user")
2 raw_train=Data("btrain")
3 raw_test=Data("btest")
4 raw_submit=Data("submit",index_col=None,header=None)
5 raw_user.shape, raw_train.shape, raw_test.shape, raw_submit.shape
```

```
/Users/qtwu/Downloads/data/users.csv
/Users/qtwu/Downloads/data/book_ratings_train.csv
/Users/qtwu/Downloads/data/book_ratings_test.csv
/Users/qtwu/Downloads/data/submission.csv
```

Out[2]:

```
((278858, 3), (260202, 3), (173469, 2), (173469, 1))
```

Build User Features By "ratings"

In [3]:

```
1 user_dict={}
2 def new_user(uid):
3     global user_dict
4     user_dict[uid]={'nbook':0,'ratings':[]}
5 for r in raw_train.to_dict('record'):
6     uid=r['User-ID']
7     if not user_dict.get(uid):
8         new_user(uid)
9     user_dict[uid]['nbook']+=1
10    user_dict[uid]['ratings'].append(r['Book-Rating'])
```

Slow caculating statistics (moment/quantile)

In [4]:

```
1 import scipy.stats
2 stat_name=['mode','tmean','tmin','tmax','skew','kurtosis']
3 def gen(uid,r):
4     result={
5         'User-ID':uid,
6         'nbook':r['nbook']
7     }
8     for name in stat_name:
9         result["rating_"+name]=eval("scipy.stats.%s(r['ratings'])"%name)
10    for v in range(9):
11        q=(v+1)*10
12        result['rating_q%d'%(v+1)]=np.percentile(r['ratings'],q)
13    return result
14 df_user_rate=pd.DataFrame([gen(uid,r) for uid,r in user_dict.items()])
15 df_user_rate['rating_mode']=df_user_rate['rating_mode'].apply(lambda r:r[0])
16 df_user_rate.sample(3)
```

Out[4]:

	User-ID	nbook	rating_kurtosis	rating_mode	rating_q1	rating_q2	rating_q3	rating_q4	rating_q5
12437	5b4f561f42	15	-0.343264	7	6.0	6.8	7.0	7.0	7.0
17102	7fcfcfb714	17	-1.603337	5	4.6	5.0	5.0	5.8	7.0
21193	d4f08d3783	4	-1.000000	5	5.9	6.8	7.7	8.2	8.0

Prepare Training Data

In [5]:

```
1 def rating_merge(rating,user): # only users.csv, books disposed
2     df=rating
3     df=df.merge(user,on='User-ID',how='left')
4     df=df.drop(['User-ID','ISBN'],1)
5     df=df.fillna(df.median()) # fill with median
6     return df
7 df_train=rating_merge(raw_train,df_user_rate)
8 df_train.sample(3)
```

Out[5]:

	Book-Rating	nbook	rating_kurtosis	rating_mode	rating_q1	rating_q2	rating_q3	rating_q4	rating_q5
243737	7	8	0.721253	9	6.1	7.0	7.1	7.8	8.0
186187	8	28	-0.668800	5	4.0	5.0	5.0	5.0	7.0
178929	9	175	-0.432006	8	7.0	7.0	7.0	8.0	8.0

Train LightGBM

(may be slow)

In [6]:

```
1 # from sklearn.ensemble import GradientBoostingClassifier, RandomForestRegressor
2 import lightgbm as lgb
3 import datetime
4
5 x_train=df_train.iloc[:,1:]
6 y_train=np.ravel(df_train.iloc[:,1])
7
8 # model=RandomForestRegressor(
9 #     500,
10 #     max_features='sqrt',
11 #     verbose=1,
12 #     n_jobs=-1,
13 # )
14
15 # model=GradientBoostingClassifier(
16 #     verbose=1,
17 #     criterion='mae',
18 #     n_estimators=10,
19 # )
20
21 model=lgb.LGBMRegressor(objective='regression')
22
23 st=datetime.datetime.now()
24 model.fit(x_train, y_train)
25 print(datetime.datetime.now()-st)
```

0:00:04.380443

Prepare Testing Data (for Submission)

In [7]:

```
1 x_test=rating_merge(raw_test,df_user_rate)
```

In [8]:

```
1 def output_test(est_name="empty"):
2     y_test_predict=model.predict(x_test)
3     y_test_predict=pd.DataFrame(y_test_predict)
4     y_test_predict.describe()
5
6     df_output=raw_submit.iloc[:,:]
7     df_output.iloc[:,0]=y_test_predict
8
9     df_output2=df_output.iloc[:,:]
10 #     df_output2=df_output2.transform(lambda x: (x/10)**1.5*10 )
11     df_output2=df_output2.round(1)
12     df_output2.to_csv("track2_%s.csv"%est_name,header=None,index=None)
13     print(df_output.describe())
14     df_output=df_output.round()
15     df_output=df_output.astype('int32')
16     df_output.to_csv("track1_%s.csv"%est_name,header=None,index=None)
17     print(df_output.describe())
18 # output_test("gbm_default_1.5down")
19
20 output_test("gbm_default_param")
21
22 # RESULT #
23 """
24 track1:
25 b04303128    2018-05-30 09:45:23
26 Banana
27 BananaBanana    1.273607
28 """
29 pass
```

```
0
count    173469.000000
mean      7.629817
std       1.071980
min       0.994106
25%       7.140688
50%       7.687896
75%       8.250696
max       9.999731
```

```
0
count    173469.000000
mean      7.683505
std       1.126445
min       1.000000
25%       7.000000
50%       8.000000
75%       8.000000
max       10.000000
```

In [10]:

```
1 # :p
2 import homework
3 from homework import *
4 reload(homework)
5 pass
```